SIMATS ENGINEERING

ASSESSMENT TOOL 3

COURSE CODE: CSA51 COURSE NAME: Cryptography and Network Security

COURSE OUTCOME COVERED

CO4: Implement best security practices for IP, email, and web service using PGP, S/MIME for enhancing email Security. (BL3, BL6)

Assessment Tool Weightage: 10%

QUESTIONS: 5 **TOTAL MARKS: 25**

Comparative Analysis

Code of Conduct:

I Avinash B (192511193) (Name / Reg No) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment.

I understand that any violation of academic integrity rules will result in disciplinary action.

Signature of the student:

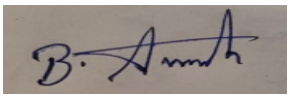

Question 1: SSL/TLS vs. IPsec for Secure Communication

Scenario: A remote employee needs (a) secure access to a company web application, and (b) secure access to the internal network as a whole. Applying both SSL/TLS and IPsec to this scenario:

Aspect	SSL/TLS	IPsec
OSI layer	Sits between Transport and Application layer	Network layer (Layer 3)
Scope of protection	Protects a single application-layer connection (e.g., one HTTPS session)	Protects all IP traffic between two endpoints/gateways, transparently to every application
Application awareness	Application (or a proxy) must be TLS-aware	Fully transparent to applications — they need no security logic of their own
Typical deployment	Per-website certificate; native browser support	Gateway/host-level configuration (IKE, security associations); used for VPNs
Granularity	Fine-grained: secures one connection/session at a time	Coarse-grained: secures an entire network path or all traffic between two hosts
Best suited for	Web-based client-server protection (HTTPS)	Network-layer protection (site-to-site VPN, protecting all traffic on a link)

Analysis: which provides better security for which scenario

For web-based protection, SSL/TLS is the better fit: it is purpose-built for securing a specific application session, is natively supported by every browser, integrates with the certificate/trust model the web already relies on, and requires no special client software beyond the browser itself.

For network-layer protection, IPsec is the better fit: because it operates below the application layer, it secures every packet between two endpoints — web traffic, email, file transfers, VoIP, anything — without requiring each individual application to implement its own security. This makes it the standard choice for site-to-site VPNs and for protecting an entire network path where transparent, blanket coverage matters more than per-application granularity.

Neither protocol is universally 'better' — they protect different layers and solve different problems. In practice, a well-designed secure communication scenario layers both: SSL/TLS to secure the specific web application session end-to-end, and an IPsec VPN to secure the overall network path the remote employee's traffic travels over, giving defense at both the application and network layers simultaneously.

Question 2: End-to-End Encryption vs. Gateway-Based Encryption for Secure Email

Scenario: Designing a secure email system for an enterprise, comparing client-side end-to-end encryption (PGP/S-MIME with user-held keys) against gateway-based encryption (applied at the organization's email gateway/perimeter).

Aspect	End-to-End Encryption	Gateway-Based Encryption
Where encryption happens	At the sender's device; decrypted only at the recipient's device	At the organizational email gateway/perimeter, on the way in/out
Confidentiality from internal IT	Strongest — no intermediate party, including the org's own mail servers/IT staff, ever sees plaintext	Weaker — plaintext is typically visible internally and at the gateway itself

Usability for end users	Requires every user to manage their own keys/certificates — complex at scale	Transparent to users; no client-side setup needed, IT manages certificates centrally
Compatibility with security scanning (AV/DLP)	Defeats gateway-based content scanning, since content is already encrypted before it leaves the device	Fully compatible — gateway can scan plaintext for malware/data-loss-prevention before encrypting outbound mail
Management overhead at scale	High — per-user key management, distribution, and recovery for potentially thousands of users	Low — centralized policy and certificate management

Justification: most suitable approach for enterprise use

For most day-to-day enterprise correspondence, gateway-based encryption is the more practical primary approach: it is transparent to users (no per-employee key management burden), keeps the organization's malware and data-loss-prevention scanning effective (since the gateway can inspect content before it is encrypted outbound), and is far easier to manage centrally across a large workforce.

However, gateway-based encryption alone leaves plaintext exposed internally and at the gateway itself, which is a meaningful weakness for especially sensitive communications (legal, HR, executive, or regulated data) where confidentiality even from internal IT staff matters, or where a compromised gateway would be catastrophic. The most suitable enterprise design is therefore a hybrid: gateway-based encryption as the default for general traffic (for manageability and scanning), with end-to-end encryption (PGP/S-MIME) offered as an option for specific high-sensitivity communications or user groups that need it — balancing usability and operational control against the strongest possible confidentiality guarantee where it is actually needed.

Question 3: HMAC vs. Digital Signatures for Data Transmission

Scenario: Two parties transmitting data need to ensure the data has not been altered and genuinely originates from the claimed sender — comparing HMAC against digital signatures for this purpose.

Aspect	HMAC	Digital Signature
Key type	Symmetric — sender and receiver share the same secret key	Asymmetric — signer uses a private key; anyone can verify with the public key
Integrity	Strong — any modification changes the MAC	Strong — any modification invalidates the signature
Authentication	Proves the message came from someone holding the shared key (which could be either party)	Proves the message came specifically from the private-key holder
Non-repudiation	Not provided — either party could have produced the same MAC, so it cannot be proven to a third party which one sent it	Provided — only the private-key holder could have produced a valid signature, verifiable by anyone
Performance	Very fast (simple hash-based computation)	Significantly slower (large modular exponentiation / elliptic-curve operations)
Key distribution	Requires securely sharing a secret key in advance	Simpler at scale — public keys can be freely distributed; only the private key must stay secret

Analysis: which provides better authentication and integrity

Both mechanisms provide strong integrity protection. For authentication, digital signatures provide the stronger guarantee: because only the private-key holder can produce a valid signature, it delivers non-repudiation and can be verified by any third party — something HMAC fundamentally cannot offer, since either party sharing the HMAC key could have generated the same tag. HMAC's authentication guarantee is therefore only as strong as the assumption that the shared key has not been exposed to anyone beyond the two intended parties.

In practice the better choice depends on the scenario: digital signatures are preferable when non-repudiation or third-party verifiability matters — legal documents, software/code signing, or the initial identity authentication in a handshake (e.g., a TLS server certificate). HMAC is preferable for high-volume, per-message integrity checking within an already-authenticated session where both parties already trust each other and share a session key — exactly how TLS itself protects the bulk data records after its handshake completes, since HMAC's much lower computational cost matters when applied to every single packet.

Question 4: Supervised vs. Unsupervised Learning for Spam Detection

Scenario: Designing a spam detection system, comparing a supervised learning approach (trained on labeled spam/ham examples) against an unsupervised approach (clustering/anomaly detection with no labels).

Aspect	Supervised Learning	Unsupervised Learning
Training data required	Labeled spam/ham examples (often abundant via user 'report spam' actions)	No labels needed — groups emails by similarity or flags statistical outliers
Strength	High accuracy on patterns similar to what it was trained on; directly optimizes the spam/ham decision	Can surface novel, previously unseen spam campaigns without needing prior labeled examples
Weakness	Needs ongoing retraining to keep up with evolving spam tactics (concept drift)	Harder to map a cluster/anomaly directly to 'spam'; generally higher false-positive risk without a ground-truth boundary
Typical role	Primary classification engine in most production filters	Complementary layer — campaign/anomaly detection alongside a supervised classifier

Evaluation: effectiveness in real-world email filtering

Supervised learning is the more directly effective approach for the core spam/ham decision, because it optimizes explicitly against real labeled outcomes and modern email providers have abundant labeled data via continuous user feedback (e.g., 'report spam' clicks), enabling frequent retraining that keeps pace with evolving spam tactics. Unsupervised learning is weaker as a standalone primary filter, since without labels it is harder to reliably distinguish 'unusual' from 'malicious', but it adds real value as a complementary layer — particularly for detecting brand-new, large-scale spam or phishing campaigns that a supervised model, trained only on historical patterns, has not yet learned to recognize.

Real-world large-scale filters (e.g., major webmail providers) do not rely on either approach in isolation: they combine a supervised classifier as the primary decision engine with unsupervised/anomaly-detection techniques and collaborative signals (many users independently reporting similar messages) as a complementary layer, achieving broader and more resilient coverage than either technique alone.

Question 5: AES vs. RSA in a Secure Communication Framework

Scenario: Designing a secure communication framework, comparing AES (symmetric) and RSA (asymmetric) encryption and determining their respective roles in achieving confidentiality and efficiency.

Aspect	AES (symmetric)	RSA (asymmetric)
Speed	Very fast — hardware-accelerated (AES-NI), suited to encrypting large volumes of data	Much slower — large modular-exponentiation operations
Data-size practicality	Efficiently handles bulk data of any size	Cannot directly encrypt data larger than its key/modulus size — impractical for bulk data
Key size for comparable security	Short keys (128–256 bits) provide strong security	Requires much larger keys (2048+ bits) for comparable security, due to a different underlying hardness assumption (factoring vs. brute-force key search)
Core strength	Efficient, strong confidentiality for the actual data	Solves the key-distribution problem — a public key can be freely shared with no pre-existing shared secret
Core weakness	Both parties must already share the same secret key — the key-distribution problem	Too slow/impractical to encrypt bulk data directly

Analysis: their roles in achieving confidentiality and efficiency

AES and RSA are complementary rather than competing choices, which is why virtually every real secure-communication framework (TLS, PGP, S/MIME, as examined in earlier assessments in this course) uses both together in a hybrid scheme. RSA is used only for the 'hard part': securely establishing or transporting a symmetric session key without requiring the two parties to already share a secret — directly solving the key-distribution problem that confidentiality depends on. AES is then used for the 'bulk part': encrypting the actual message or file data efficiently, since it achieves strong confidentiality at a small fraction of RSA's computational cost per byte processed.

Neither algorithm alone is sufficient for a practical framework: pure RSA is too slow and, for anything larger than its modulus, outright unable to encrypt bulk data directly; pure AES alone re-introduces the very key-distribution problem RSA is meant to solve, since the two parties would need some other secure channel to agree on the AES key in the first place. The hybrid architecture — RSA (or another asymmetric algorithm) for key exchange and authentication, AES for bulk data confidentiality — is therefore the standard, practical design that achieves both strong confidentiality and real-world efficiency simultaneously.